

Ashish Shandil

Data Engineer – Cloud Data Pipelines, ETL/ELT & AI-Enabled Data Infrastructure

600-660-3281 | ashishshandil27@gmail.com | [LinkedIn](#) | [GitHub](#)

Professional Summary

Data Engineer with 3+ years designing, building, and maintaining scalable ETL/ELT pipelines, cloud data integrations, and analytical platforms – having built and optimized 30+ data pipelines and 15+ Power BI dashboards, processing 500GB+ of data across enterprise systems. Skilled in implementing data quality validation, monitoring, and observability frameworks, optimizing SQL/Spark workloads, and delivering CI/CD-driven deployments alongside engineering and analytics stakeholders. Databricks Certified Data Engineer Associate and Snowflake SnowPro Core certified, with hands-on experience building retrieval-augmented generation (RAG) AI agents and ML models – bringing a strong foundation for AI-enabled data infrastructure, knowledge grounding, and modern cloud data platforms.

Key Skills

- **Languages & Data Processing:** Python, SQL, PySpark, Pandas
- **Data Pipelines & Integration:** ETL/ELT Design, Data Ingestion Pipelines, System Integrations, Azure Data Factory, Azure Databricks, Microsoft Fabric, Delta Lake
- **Data Architecture:** Data Modeling, Schema Design, Query Optimization, Data Validation & Governance, Relational & NoSQL Databases
- **Cloud & Platforms:** Microsoft Azure (Databricks, Synapse, Data Lake), Snowflake (SnowPro Core Certified)
- **DevOps & Reliability:** Git, GitHub, Azure DevOps CI/CD, Automated Monitoring & Validation
- **AI/ML & Emerging Data:** Retrieval-Augmented Generation (RAG), AI Agent Development (Microsoft Copilot Studio), Knowledge Grounding, Vector-Based Search, XGBoost, Supervised Learning
- **Analytics & BI:** Power BI, DAX, KPI Dashboards, Semantic Modeling

Professional Experience

MAQ Software, Noida, India — *Data Engineer 2* (Nov 2025 – Present)

- Design and maintain scalable ETL/ELT pipelines using Microsoft Fabric, Azure Databricks, Synapse Analytics, and Azure Data Factory, monitoring and reviewing 25+ pipelines across production workloads.
- Implement automated monitoring frameworks with configurable timeouts, retry intervals, and AI-driven alerting to proactively detect pipeline failures, and automate stakeholder reporting on pipeline execution status via scheduled email and Microsoft Teams notifications.
- Optimize Spark workloads and SQL transformations across 500GB+ of enterprise data, improving pipeline execution efficiency and reporting performance for enterprise stakeholders.
- Build and deploy RAG-based AI agents across 2 Microsoft Copilot Studio environments and multiple GitHub Copilot-based agents, automating internal project work-item workflows and grounding generative responses in enterprise knowledge sources via vector-based retrieval; also developing a supervised machine learning model (XGBoost, approx. 85% accuracy in testing) applying feature engineering and evaluation best practices.

MAQ Software, Noida, India — *Data Engineer 1* (Aug 2024 – Nov 2025)

- Built and optimized 30+ production data pipelines using Azure Databricks, Microsoft Fabric, Synapse Analytics, and SQL, supporting enterprise-scale data integration and reporting needs.
- Led migration of legacy data infrastructure to Microsoft Fabric, improving platform scalability, maintainability, and long-term reliability.
- Automated ETL workflows and data quality checks; developed and enhanced 15+ Power BI dashboards with DAX-based semantic models for enterprise KPI reporting.
- Partnered with analytics and business stakeholders to gather requirements and deliver scalable, well-governed data solutions using Delta Lake architecture.

MAQ Software, Hyderabad, India — *Associate Software Engineer* (Aug 2023 – Aug 2024)

- Contributed to ETL pipeline development, data validation, and reporting solutions using Azure Data Factory, Databricks, and SQL, improving reporting efficiency by 20%.
- Performed data quality and KPI validation testing, resolving 10+ production issues and reducing reporting discrepancies by 30% in collaboration with senior engineers.

Projects

Code Analyzer Automation Tool – Built a Python/SQL automation tool to analyze repository files, clean metadata, and load processed data into data warehouse tables, reducing manual analysis effort from 3–4 story points to under 1 story point per task (approx. 75% reduction).

ML Prediction Model – XGBoost – Developing a supervised learning model for a predictive analytics use case, currently achieving approx. 85% accuracy in testing; applied feature engineering and evaluation metrics.

Education

Chitkara University, Solan, Himachal Pradesh — Bachelor of Engineering in Computer Science (Aug 2020 – May 2024)

Certifications

- Databricks Certified Data Engineer Associate
- Snowflake SnowPro Core Certification (Score: 913)
- Microsoft Fabric DP-700 | Microsoft Fabric DP-600 (Score: 913)
- Microsoft SQL DP-800 (Score: 905) | Microsoft Azure Developer Associate (AZ-204)
- Microsoft Agentic AI Business Solutions Architect AB-100 (Score: 767)